

# CPEgenerator v2 + KGCS — pilot status

MOTHERHACKER · R&D · OPEN-SOURCE

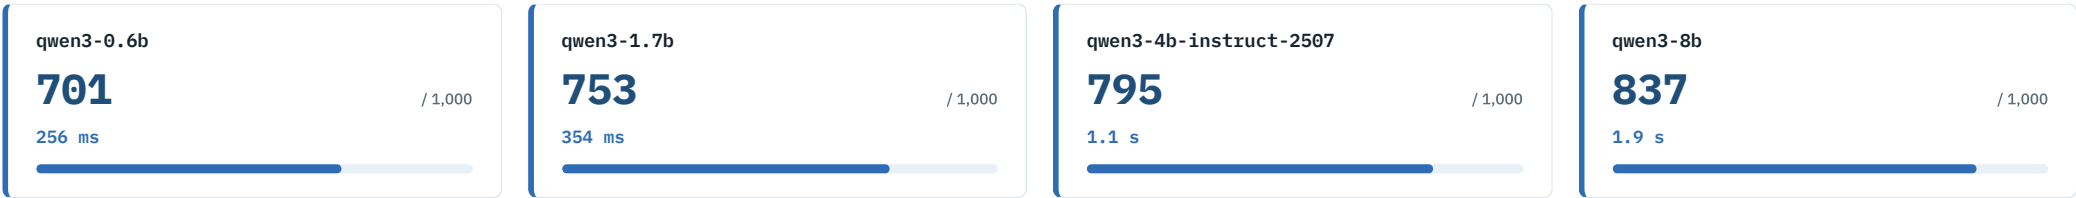
## High-confidence automatic CPE resolution (M1x)



2023 figure measured on raw inventory; 2026 figure on clean annotated titles (exact CPE on gold-1k: **83.7%**). Raw-inventory mass run in progress (cascade qwen3-1.7b → qwen3-8b over 487k titles); early real-inventory signal: **20% M1x** on a live 82-title Windows inventory vs the 4.9% baseline.

## Local quality / cost curve

Exact CPEs per 1,000 gold titles @ median latency per title — RTX 3060 12GB, LM Studio, temperature 0, reasoning off



- METHODThe LLM only proposes entities as JSON; a deterministic ABNF validator and the M1–M3 matcher decide. Dictionary: 1.77M CPEs served locally from the KGCS Neo4j graph. 177 offline tests.
- STATUS14 local models benchmarked; extraction-mode decision closed (single-call beats per-field at every size); two-model cascade selected for the 487k raw run.